

Add a New System Call to list all processes

Target

1. Add a new system call into the linux kernel
2. Test the new system call in user model

Tools

Install GCC Software Colletion

```
sudo apt-get install build-essential
```

How to use GCC

- gcc and make

How to do

see the pdf document: newsyscall_step2.pdf

Step0

new customized kernel config: config1, config2

Step1 (Linux kernel 6.5.7)

include/linux/syscalls.h

在文件(No. 1176)

#endif /* CONFIG_ARCH_HAS_SYSCALL_WRAPPER */之前, 添加一行:

```
asmlinkage long sys_schello(void);
```

Step2 (Linux kernel 6.5.7)

kernel/sys.c

在文件SYSCALL_DEFINE0(gettid)函数之后 (No. 958) , 添加如下行:

```
SYSCALL_DEFINE0(schello)
{
    printk("Hello new system call schello!Your ID\n");
    return 0;
}
```

Step3 (Linux kernel 6.5.7)

针对64位OS

arch/x86/entry/syscalls/syscall_64.tbl

在 334 common resq sys_rseq 行之后, 添加如下行:

```
335 common schello sys_schello
```

Step4

```
make clean
```

```
make -j5
```

```
sudo make modules_install
```

```
sudo make install
```

Step 5

重新启动:

```
reboot
```

确认新内核是否成功运行:

```
uname -a
```

```
hxusr@hxhost:~$ uname -a
Linux hxhost 6.5.72110731 #2 SMP PREEMPT_DYNAMIC Thu Oct 26 03:14:20 CST 2023 x86_64 x86_64 x86_64 GNU/Linux
```

Step 6

编写用户态测试程序testschello.c

```
#include <unistd.h>
#include <sys/syscall.h>
#include <sys/types.h>
#include <stdio.h>
#define _NR_schello 335
int main(int argc, char *argv[])
{
    syscall(_NR_schello);
    print("ok! run dmesg | grep hello in terminal!\n");
    return 0;
}
```

Step 7

编译用户态测试程序testschello.c, 并执行

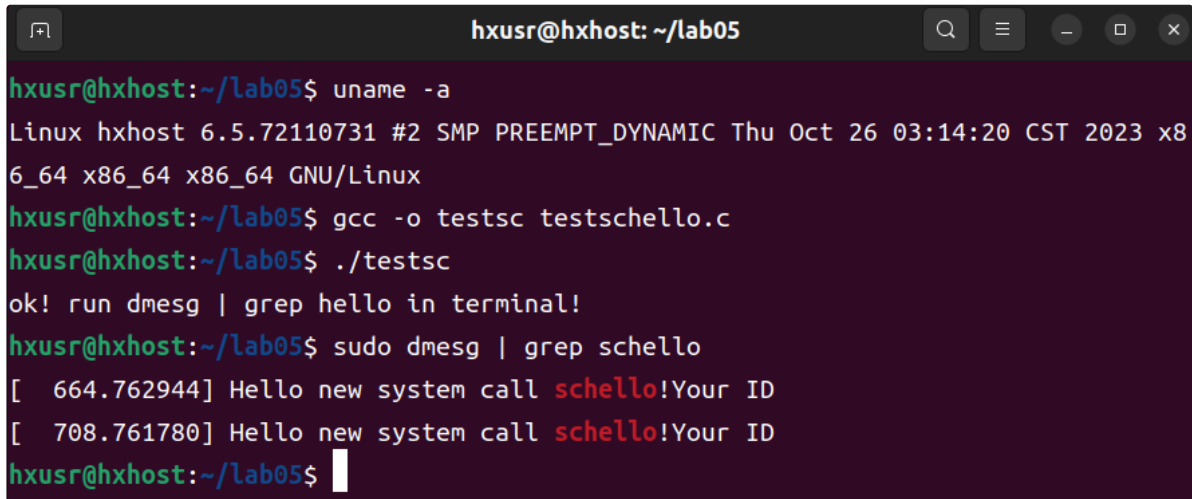
```
gcc -o testsc testschello.c

$ sudo dmesg -C

./testsc

$dmesg | grep schello

[ 1648.215250] Hello new system call schello!
```

A terminal window titled 'hxusr@hxhost: ~/lab05' with standard window controls. The terminal shows the following commands and output:

```
hxusr@hxhost:~/lab05$ uname -a
Linux hxhost 6.5.72110731 #2 SMP PREEMPT_DYNAMIC Thu Oct 26 03:14:20 CST 2023 x86_64 x86_64 x86_64 GNU/Linux
hxusr@hxhost:~/lab05$ gcc -o testsc testschello.c
hxusr@hxhost:~/lab05$ ./testsc
ok! run dmesg | grep hello in terminal!
hxusr@hxhost:~/lab05$ sudo dmesg | grep schello
[ 664.762944] Hello new system call schello!Your ID
[ 708.761780] Hello new system call schello!Your ID
hxusr@hxhost:~/lab05$
```

Add new system call to list all processes

实验要求

需要列出每个进程的名字(comm)、进程ID号(pid)、父进程ID号、进程状态、学号姓名等

需要列出每个进程的名字(comm)、进程ID号(pid)、父进程ID号、进程状态、学号姓名等

需要列出每个进程的名字(comm)、进程ID号(pid)、父进程ID号、进程状态、学号姓名等

参考内容

进程结构体:struct task_struct *p;

双向链表:

```
struct list_head {
    struct list_head *next, *prev;
};
```

第一个进程:init_task

当前的集成:current

进程的兄弟进程:struct list_head tasks;

进程p的下一个进程:next_task(p)

```
#define next_task(p) \
    list_entry_rcu((p)→tasks.next, struct task_struct, tasks)
```

遍历所有进程:

```
#define for_each_process(p) \
    for (p = &init_task ; (p = next_task(p)) != &init_task ; )
```

进程的状态:static inline char task_state_to_char(struct task_struct *tsk)

实现方式

在上述schello基础上修改kernel/sys.c文件

示例：在文件SYSCALL_DEFINE0(gettid)函数之后，添加如下行：

```
SYSCALL_DEFINE0(schello)
{
    struct task_struct *p;
    printf("Hello new system call schello!\n");
    printk("%-20s %-6s %-6s\n", "Name", "Pid", "Stat");
    for (p = &init_task; (p = next_task(p)) != &init_task;)
        printk("%-20s %-6d %-6ld\n", p->comm, p->pid, p->state);
    return 0;
}
```

请注意实验要求：需要列出每个进程的名字(comm)、进程ID号(pid)、父进程ID号、进程状态、学号姓名等

实验结果图

sudo dmesg查看log

```
[ 220.108934] Hello new system call schello!Your ID
[ 220.108937] Hello new system call schello!Your ID178
[ 220.108938] print Name p_Pid Pid Stat
[ 220.108939] systemd 0 1 1
[ 220.108941] kthreadd 0 2 1
[ 220.108942] rcu_gp 2 3 1026
[ 220.108943] rcu_par_gp 2 4 1026
[ 220.108944] slub_flushwq 2 5 1026
[ 220.108945] netns 2 6 1026
[ 220.108946] kworker/0:0 2 7 1026
[ 220.108946] kworker/0:0H 2 8 1026
[ 220.108947] kworker/0:1 2 9 1026
[ 220.108948] kworker/u4:0 2 10 1026
[ 220.108949] mm_percpu_wq 2 11 1026
[ 220.108950] rcu_tasks_kthre 2 12 1026
[ 220.108951] rcu_tasks_rude_ 2 13 1026
[ 220.108952] rcu_tasks_trace 2 14 1026
[ 220.108953] ksoftirqd/0 2 15 1
[ 220.108954] rcu_preempt 2 16 1026
[ 220.108955] migration/0 2 17 1
[ 220.108956] idle_inject/0 2 18 1
[ 220.108957] cpuhp/0 2 19 1
[ 220.108958] cpuhp/1 2 20 1
[ 220.108959] idle_inject/1 2 21 1
```

process in Linux kernel

涉及文件:

`include/linux/sched.h`

```
struct task_struct {
#ifdef CONFIG_THREAD_INFO_IN_TASK
    /*
     * For reasons of header soup (see current_thread_info()), this
     * must be the first element of task_struct.
     */
    struct thread_info    thread_info;
#endif
    unsigned int          __state;

#ifdef CONFIG_PREEMPT_RT
    /* saved state for "spinlock sleepers" */
    unsigned int          saved_state;
#endif

    /*
     * This begins the randomizable portion of task_struct. Only
     * scheduling-critical items should be added above here.
     */
    randomized_struct_fields_start

    void                  *stack;
    refcount_t            usage;
    /* Per task flags (PF_*), defined further below: */
    unsigned int          flags;
    unsigned int          ptrace;

#ifdef CONFIG_SMP
    int                   on_cpu;
    struct __call_single_node wake_entry;
    unsigned int          wakee_flips;
    unsigned long         wakee_flip_decay_ts;
    struct task_struct    *last_wakee;

    /*
     * recent_used_cpu is initially set as the last CPU used by a task
     * that wakes affine another task. Waker/wakee relationships can
     * push tasks around a CPU where each wakeup moves to the next one.
     * Tracking a recently used CPU allows a quick search for a recently
     * used CPU that may be idle.
     */
    int                   recent_used_cpu;
    int                   wake_cpu;
#endif
    int                   on_rq;

    int                   prio;
    int                   static_prio;
```

```

int          normal_prio;
unsigned int  rt_priority;

struct sched_entity se;
struct sched_rt_entity rt;
struct sched_dl_entity dl;
const struct sched_class *sched_class;

#ifdef CONFIG_SCHED_CORE
struct rb_node core_node;
unsigned long core_cookie;
unsigned int core_occupation;
#endif

#ifdef CONFIG_CGROUP_SCHED
struct task_group *sched_task_group;
#endif

#ifdef CONFIG_UCLAMP_TASK
/*
 * Clamp values requested for a scheduling entity.
 * Must be updated with task_rq_lock() held.
 */
struct uclamp_se uclamp_req[UCLAMP_CNT];
/*
 * Effective clamp values used for a scheduling entity.
 * Must be updated with task_rq_lock() held.
 */
struct uclamp_se uclamp[UCLAMP_CNT];
#endif

struct sched_statistics stats;

#ifdef CONFIG_PREEMPT_NOTIFIERS
/* List of struct preempt_notifier: */
struct hlist_head preempt_notifiers;
#endif

#ifdef CONFIG_BLK_DEV_IO_TRACE
unsigned int btrace_seq;
#endif

unsigned int policy;
int nr_cpus_allowed;
const cpumask_t *cpus_ptr;
cpumask_t *user_cpus_ptr;
cpumask_t cpus_mask;
void *migration_pending;
#ifdef CONFIG_SMP
unsigned short migration_disabled;
#endif
unsigned short migration_flags;

#ifdef CONFIG_PREEMPT_RCU

```

```

int            rcu_read_lock_nesting;
union rcu_special      rcu_read_unlock_special;
struct list_head      rcu_node_entry;
struct rcu_node        *rcu_blocked_node;
#endif /* #ifdef CONFIG_PREEMPT_RCU */

#ifdef CONFIG_TASKS_RCU
    unsigned long      rcu_tasks_nvcsw;
    u8                  rcu_tasks_holdout;
    u8                  rcu_tasks_idx;
    int                 rcu_tasks_idle_cpu;
    struct list_head     rcu_tasks_holdout_list;
#endif /* #ifdef CONFIG_TASKS_RCU */

#ifdef CONFIG_TASKS_TRACE_RCU
    int                 trc_reader_nesting;
    int                 trc_ipi_to_cpu;
    union rcu_special     trc_reader_special;
    struct list_head      trc_holdout_list;
    struct list_head      trc_blkdn_node;
    int                 trc_blkdn_cpu;
#endif /* #ifdef CONFIG_TASKS_TRACE_RCU */

    struct sched_info     sched_info;

    struct list_head      tasks;
#ifdef CONFIG_SMP
    struct plist_node     pushable_tasks;
    struct rb_node         pushable_dl_tasks;
#endif

    struct mm_struct      *mm;
    struct mm_struct      *active_mm;

    int                   exit_state;
    int                   exit_code;
    int                   exit_signal;
    /* The signal sent when the parent dies: */
    int                   pdeath_signal;
    /* JOBCTL_*, siglock protected: */
    unsigned long         jobctl;

    /* Used for emulating ABI behavior of previous Linux versions: */
    unsigned int          personality;

    /* Scheduler bits, serialized by scheduler locks: */
    unsigned              sched_reset_on_fork:1;
    unsigned              sched_contributes_to_load:1;
    unsigned              sched_migrated:1;

    /* Force alignment to the next boundary: */
    unsigned              :0;

    /* Unserialized, strictly 'current' */

```

```

/*
 * This field must not be in the scheduler word above due to wakelist
 * queueing no longer being serialized by p->on_cpu. However:
 *
 * p->XXX = X;          ttwu()
 * schedule()           if (p->on_rq && ..) // false
 * smp_mb__after_spinlock(); if (smp_load_acquire(&p->on_cpu) &&
//true
 * deactivate_task()      ttwu_queue_wakelist()
 * p->on_rq = 0;           p->sched_remote_wakeup = Y;
 *
 * guarantees all stores of 'current' are visible before
 * ->sched_remote_wakeup gets used, so it can be in this word.
 */
unsigned          sched_remote_wakeup:1;

/* Bit to tell LSMs we're in execve(): */
unsigned          in_execve:1;
unsigned          in_iowait:1;
#ifdef TIF_RESTORE_SIGMASK
unsigned          restore_sigmask:1;
#endif
#ifdef CONFIG_MEMCG
unsigned          in_user_fault:1;
#endif
#ifdef CONFIG_LRU_GEN
/* whether the LRU algorithm may apply to this access */
unsigned          in_lru_fault:1;
#endif
#ifdef CONFIG_COMPAT_BRK
unsigned          brk_randomized:1;
#endif
#ifdef CONFIG_CGROUPS
/* disallow userland-initiated cgroup migration */
unsigned          no_cgroup_migration:1;
/* task is frozen/stopped (used by the cgroup freezer) */
unsigned          frozen:1;
#endif
#ifdef CONFIG_BLK_CGROUP
unsigned          use_memdelay:1;
#endif
#ifdef CONFIG_PSI
/* Stalled due to lack of memory */
unsigned          in_memstall:1;
#endif
#ifdef CONFIG_PAGE_OWNER
/* Used by page_owner=on to detect recursion in page tracking. */
unsigned          in_page_owner:1;
#endif
#ifdef CONFIG_EVENTFD
/* Recursion prevention for eventfd_signal() */
unsigned          in_eventfd:1;
#endif

```



```

#ifdef CONFIG_IOMMU_SVA
    unsigned          pasid_activated:1;
#endif
#ifdef CONFIG_CPU_SUP_INTEL
    unsigned          reported_split_lock:1;
#endif
#ifdef CONFIG_TASK_DELAY_ACCT
    /* delay due to memory thrashing */
    unsigned          in_thrashing:1;
#endif

    unsigned long      atomic_flags; /* Flags requiring atomic access. */

    struct restart_block restart_block;

    pid_t              pid;
    pid_t              tgid;

#ifdef CONFIG_STACKPROTECTOR
    /* Canary value for the -fstack-protector GCC feature: */
    unsigned long      stack_canary;
#endif
/*
 * Pointers to the (original) parent process, youngest child, younger
 sibling,
 * older sibling, respectively. (p→father can be replaced with
 * p→real_parent→pid)
 */

/* Real parent process: */
struct task_struct __rcu *real_parent;

/* Recipient of SIGCHLD, wait4() reports: */
struct task_struct __rcu *parent;

/*
 * Children/sibling form the list of natural children:
 */
struct list_head      children;
struct list_head      sibling;
struct task_struct     *group_leader;

/*
 * 'ptraced' is the list of tasks this task is using ptrace() on.
 *
 * This includes both natural children and PTRACE_ATTACH targets.
 * 'ptrace_entry' is this task's link on the p→parent→ptraced list.
 */
struct list_head      ptraced;
struct list_head      ptrace_entry;

/* PID/PID hash table linkage. */
struct pid            *thread_pid;
struct hlist_node      pid_links[PIDTYPE_MAX];

```

```

    struct list_head      thread_group;
    struct list_head      thread_node;

    struct completion      *vfork_done;

    /* CLONE_CHILD_SETTID: */
    int __user              *set_child_tid;

    /* CLONE_CHILD_CLEARTID: */
    int __user              *clear_child_tid;

    /* PF_KTHREAD | PF_IO_WORKER */
    void                    *worker_private;

    u64                     utime;
    u64                     stime;
#ifdef CONFIG_ARCH_HAS_SCALED_CPUTIME
    u64                     utimescaled;
    u64                     stimescaled;
#endif
    u64                     gtime;
    struct prev_cputime      prev_cputime;
#ifdef CONFIG_VIRT_CPU_ACCOUNTING_GEN
    struct vtime             vtime;
#endif

#ifdef CONFIG_NO_HZ_FULL
    atomic_t                tick_dep_mask;
#endif
    /* Context switch counts: */
    unsigned long            nvcs;
    unsigned long            nivcs;

    /* Monotonic time in nsecs: */
    u64                     start_time;

    /* Boot based time in nsecs: */
    u64                     start_boottime;

    /* MM fault and swap info: this can arguably be seen as either mm-specific
or thread-specific: */
    unsigned long            minflt;
    unsigned long            majflt;

    /* Empty if CONFIG_POSIX_CPUTIMERS=n */
    struct posix_cputimers    posix_cputimers;

#ifdef CONFIG_POSIX_CPU_TIMERS_TASK_WORK
    struct posix_cputimers_work posix_cputimers_work;
#endif

    /* Process credentials: */

    /* Tracer's credentials at attach: */

```

```

const struct cred __rcu      *ptracer_cred;

/* Objective and real subjective task credentials (COW): */
const struct cred __rcu      *real_cred;

/* Effective (overridable) subjective task credentials (COW): */
const struct cred __rcu      *cred;

#ifdef CONFIG_KEYS
/* Cached requested key. */
struct key                    *cached_requested_key;
#endif

/*
 * executable name, excluding path.
 *
 * - normally initialized setup_new_exec()
 * - access it with [gs]et_task_comm()
 * - lock it with task_lock()
 */
char                            comm[TASK_COMM_LEN];

struct nameidata                *nameidata;

#ifdef CONFIG_SYSVIPC
struct sysv_sem                 sysvsem;
struct sysv_shm                 sysvshm;
#endif
#ifdef CONFIG_DETECT_HUNG_TASK
unsigned long                    last_switch_count;
unsigned long                    last_switch_time;
#endif
/* Filesystem information: */
struct fs_struct                *fs;

/* Open file information: */
struct files_struct              *files;

#ifdef CONFIG_IO_URING
struct io_uring_task            *io_uring;
#endif

/* Namespaces: */
struct nsproxy                  *nsproxy;

/* Signal handlers: */
struct signal_struct              *signal;
struct sighand_struct __rcu      *sighand;
sigset_t                        blocked;
sigset_t                        real_blocked;
/* Restored if set_restore_sigmask() was used: */
sigset_t                        saved_sigmask;
struct sigpending                pending;
unsigned long                    sas_ss_sp;

```

```

size_t          sas_ss_size;
unsigned int     sas_ss_flags;

struct callback_head *task_works;

#ifdef CONFIG_AUDIT
#ifdef CONFIG_AUDITSYSCALL
    struct audit_context *audit_context;
#endif
#endif
    kuid_t          loginuid;
    unsigned int     sessionid;
#endif
    struct seccomp     seccomp;
    struct syscall_user_dispatch syscall_dispatch;

    /* Thread group tracking: */
    u64          parent_exec_id;
    u64          self_exec_id;

    /* Protection against (de-)allocation: mm, files, fs, tty, keyrings,
    mems_allowed, mempolicy: */
    spinlock_t     alloc_lock;

    /* Protection of the PI data structures: */
    raw_spinlock_t pi_lock;

    struct wake_q_node wake_q;

#ifdef CONFIG_RT_MUTEXES
    /* PI waiters blocked on a rt_mutex held by this task: */
    struct rb_root_cached pi_waiters;
    /* Updated under owner's pi_lock and rq lock */
    struct task_struct *pi_top_task;
    /* Deadlock detection and priority inheritance handling: */
    struct rt_mutex_waiter *pi_blocked_on;
#endif

#ifdef CONFIG_DEBUG_MUTEXES
    /* Mutex deadlock detection: */
    struct mutex_waiter *blocked_on;
#endif

#ifdef CONFIG_DEBUG_ATOMIC_SLEEP
    int non_block_count;
#endif

#ifdef CONFIG_TRACE_IRQFLAGS
    struct irqtrace_events irqtrace;
    unsigned int          hardirq_threaded;
    u64          hardirq_chain_key;
    int          softirqs_enabled;
    int          softirq_context;
    int          irq_config;
#endif

```

```

#ifdef CONFIG_PREEMPT_RT
    int                softirq_disable_cnt;
#endif

#ifdef CONFIG_LOCKDEP
# define MAX_LOCK_DEPTH      48UL
    u64                curr_chain_key;
    int                lockdep_depth;
    unsigned int        lockdep_recursion;
    struct held_lock     held_locks[MAX_LOCK_DEPTH];
#endif

#if defined(CONFIG_UBSAN) && !defined(CONFIG_UBSAN_TRAP)
    unsigned int        in_ubsan;
#endif

    /* Journalling filesystem info: */
    void                *journal_info;

    /* Stacked block device info: */
    struct bio_list      *bio_list;

    /* Stack plugging: */
    struct blk_plug      *plug;

    /* VM state: */
    struct reclaim_state *reclaim_state;

    struct io_context    *io_context;

#ifdef CONFIG_COMPACTION
    struct capture_control *capture_control;
#endif

    /* Ptrace state: */
    unsigned long        ptrace_message;
    kernel_siginfo_t     *last_siginfo;

    struct task_io_accounting ioac;
#ifdef CONFIG_PSI
    /* Pressure stall state */
    unsigned int         psi_flags;
#endif
#ifdef CONFIG_TASK_XACCT
    /* Accumulated RSS usage: */
    u64                  acct_rss_mem1;
    /* Accumulated virtual memory usage: */
    u64                  acct_vm_mem1;
    /* stime + utime since last update: */
    u64                  acct_timexpd;
#endif
#ifdef CONFIG_CPUSETS
    /* Protected by →alloc_lock: */
    nodemask_t           mems_allowed;
    /* Sequence number to catch updates: */

```

```

    seqcount_spinlock_t    mems_allowed_seq;
    int                    cpuset_mem_spread_rotor;
    int                    cpuset_slab_spread_rotor;
#endif
#ifdef CONFIG_CGROUPS
    /* Control Group info protected by css_set_lock: */
    struct css_set __rcu    *cgroups;
    /* cg_list protected by css_set_lock and tsk->alloc_lock: */
    struct list_head        cg_list;
#endif
#ifdef CONFIG_X86_CPU_RESCTRL
    u32                    closid;
    u32                    rmid;
#endif
#ifdef CONFIG_FUTEX
    struct robust_list_head __user *robust_list;
#endif
#ifdef CONFIG_COMPAT
    struct compat_robust_list_head __user *compat_robust_list;
#endif
    struct list_head        pi_state_list;
    struct futex_pi_state    *pi_state_cache;
    struct mutex            futex_exit_mutex;
    unsigned int            futex_state;
#endif
#ifdef CONFIG_PERF_EVENTS
    struct perf_event_context *perf_event_ctxp;
    struct mutex            perf_event_mutex;
    struct list_head        perf_event_list;
#endif
#ifdef CONFIG_DEBUG_PREEMPT
    unsigned long            preempt_disable_ip;
#endif
#ifdef CONFIG_NUMA
    /* Protected by alloc_lock: */
    struct mempolicy        *mempolicy;
    short                    il_prev;
    short                    pref_node_fork;
#endif
#ifdef CONFIG_NUMA_BALANCING
    int                    numa_scan_seq;
    unsigned int            numa_scan_period;
    unsigned int            numa_scan_period_max;
    int                    numa_preferred_nid;
    unsigned long            numa_migrate_retry;
    /* Migration stamp: */
    u64                    node_stamp;
    u64                    last_task_numa_placement;
    u64                    last_sum_exec_runtime;
    struct callback_head    numa_work;

    /*
     * This pointer is only modified for current in syscall and
     * pagefault context (and for tasks being destroyed), so it can be read
     * from any of the following contexts:

```

```

* - RCU read-side critical section
* - current→numa_group from everywhere
* - task's runqueue locked, task not running
*/
struct numa_group __rcu      *numa_group;

/*
 * numa_faults is an array split into four regions:
 * faults_memory, faults_cpu, faults_memory_buffer, faults_cpu_buffer
 * in this precise order.
 *
 * faults_memory: Exponential decaying average of faults on a per-node
 * basis. Scheduling placement decisions are made based on these
 * counts. The values remain static for the duration of a PTE scan.
 * faults_cpu: Track the nodes the process was running on when a NUMA
 * hinting fault was incurred.
 * faults_memory_buffer and faults_cpu_buffer: Record faults per node
 * during the current scan window. When the scan completes, the counts
 * in faults_memory and faults_cpu decay and these values are copied.
 */
unsigned long      *numa_faults;
unsigned long      total_numa_faults;

/*
 * numa_faults_locality tracks if faults recorded during the last
 * scan window were remote/local or failed to migrate. The task scan
 * period is adapted based on the locality of the faults with different
 * weights depending on whether they were shared or private faults
 */
unsigned long      numa_faults_locality[3];

unsigned long      numa_pages_migrated;
#endif /* CONFIG_NUMA_BALANCING */

#ifdef CONFIG_RSEQ
struct rseq __user *rseq;
u32 rseq_len;
u32 rseq_sig;
/*
 * RmW on rseq_event_mask must be performed atomically
 * with respect to preemption.
 */
unsigned long rseq_event_mask;
#endif

#ifdef CONFIG_SCHED_MM_CID
int      mm_cid;      /* Current cid in mm */
int      last_mm_cid; /* Most recent cid in mm */
int      migrate_from_cpu;
int      mm_cid_active; /* Whether cid bitmap is active */
struct callback_head      cid_work;
#endif

struct tlbflush_unmap_batch tlb_ubic;

```

```

/* Cache last used pipe for splice(): */
struct pipe_inode_info      *splice_pipe;

struct page_frag            task_frag;

#ifdef CONFIG_TASK_DELAY_ACCT
    struct task_delay_info    *delays;
#endif

#ifdef CONFIG_FAULT_INJECTION
    int                        make_it_fail;
    unsigned int               fail_nth;
#endif

/*
 * When (nr_dirtied ≥ nr_dirtied_pause), it's time to call
 * balance_dirty_pages() for a dirty throttling pause:
 */
int                            nr_dirtied;
int                            nr_dirtied_pause;
/* Start of a write-and-pause period: */
unsigned long                  dirty_paused_when;

#ifdef CONFIG_LATENCYTOP
    int                        latency_record_count;
    struct latency_record      latency_record[LT_SAVECOUNT];
#endif

/*
 * Time slack values; these are used to round up poll() and
 * select() etc timeout values. These are in nanoseconds.
 */
u64                            timer_slack_ns;
u64                            default_timer_slack_ns;

#if defined(CONFIG_KASAN_GENERIC) || defined(CONFIG_KASAN_SW_TAGS)
    unsigned int               kasan_depth;
#endif

#ifdef CONFIG_KCSAN
    struct kcsan_ctx            kcsan_ctx;
#endif
#ifdef CONFIG_TRACE_IRQFLAGS
    struct irqtrace_events      kcsan_save_irqtrace;
#endif
#ifdef CONFIG_KCSAN_WEAK_MEMORY
    int                         kcsan_stack_depth;
#endif

#ifdef CONFIG_KMSAN
    struct kmsan_ctx            kmsan_ctx;
#endif

#if IS_ENABLED(CONFIG_KUNIT)
    struct kunit                *kunit_test;

```



```

#endif

#ifdef CONFIG_FUNCTION_GRAPH_TRACER
    /* Index of current stored address in ret_stack: */
    int curr_ret_stack;
    int curr_ret_depth;

    /* Stack of return addresses for return function tracing: */
    struct ftrace_ret_stack *ret_stack;

    /* Timestamp for last schedule: */
    unsigned long long ftrace_timestamp;

    /*
     * Number of functions that haven't been traced
     * because of depth overrun:
     */
    atomic_t trace_overrun;

    /* Pause tracing: */
    atomic_t tracing_graph_pause;
#endif

#ifdef CONFIG_TRACING
    /* Bitmask and counter of trace recursion: */
    unsigned long trace_recursion;
#endif /* CONFIG_TRACING */

#ifdef CONFIG_KCOV
    /* See kernel/kcov.c for more details. */

    /* Coverage collection mode enabled for this task (0 if disabled): */
    unsigned int kcov_mode;

    /* Size of the kcov_area: */
    unsigned int kcov_size;

    /* Buffer for coverage collection: */
    void *kcov_area;

    /* KCOV descriptor wired with this task or NULL: */
    struct kcov *kcov;

    /* KCOV common handle for remote coverage collection: */
    u64 kcov_handle;

    /* KCOV sequence number: */
    int kcov_sequence;

    /* Collect coverage from softirq context: */
    unsigned int kcov_softirq;
#endif

#ifdef CONFIG_MEMCG

```

```

    struct mem_cgroup      *memcg_in_oom;
    gfp_t                  memcg_oom_gfp_mask;
    int                    memcg_oom_order;

    /* Number of pages to reclaim on returning to userland: */
    unsigned int           memcg_nr_pages_over_high;

    /* Used by memcontrol for targeted memcg charge: */
    struct mem_cgroup      *active_memcg;
#endif

#ifdef CONFIG_BLK_CGROUP
    struct gendisk          *throttle_disk;
#endif

#ifdef CONFIG_UPROBES
    struct uprobe_task      *utask;
#endif
#if defined(CONFIG_BCACHE) || defined(CONFIG_BCACHE_MODULE)
    unsigned int            sequential_io;
    unsigned int            sequential_io_avg;
#endif
    struct kmap_ctrl        kmap_ctrl;
#ifdef CONFIG_DEBUG_ATOMIC_SLEEP
    unsigned long           task_state_change;
# ifdef CONFIG_PREEMPT_RT
    unsigned long           saved_state_change;
# endif
#endif
    struct rcu_head         rcu;
    refcount_t              rcu_users;
    int                     pagefault_disabled;
#ifdef CONFIG_MMU
    struct task_struct      *oom_reaper_list;
    struct timer_list       oom_reaper_timer;
#endif
#ifdef CONFIG_VMAP_STACK
    struct vm_struct        *stack_vm_area;
#endif
#ifdef CONFIG_THREAD_INFO_IN_TASK
    /* A live task holds one reference: */
    refcount_t              stack_refcount;
#endif
#ifdef CONFIG_LIVEPATCH
    int                     patch_state;
#endif
#ifdef CONFIG_SECURITY
    /* Used by LSM modules for access restriction: */
    void                    *security;
#endif
#ifdef CONFIG_BPF_SYSCALL
    /* Used by BPF task local storage */
    struct bpf_local_storage __rcu *bpf_storage;
    /* Used for BPF run context */

```

```

    struct bpf_run_ctx      *bpf_ctx;
#endif

#ifdef CONFIG_GCC_PLUGIN_STACKLEAK
    unsigned long           lowest_stack;
    unsigned long           prev_lowest_stack;
#endif

#ifdef CONFIG_X86_MCE
    void __user             *mce_vaddr;
    __u64                   mce_kflags;
    u64                     mce_addr;
    __u64                   mce_ripv : 1,
                           mce_whole_page : 1,
                           __mce_reserved : 62;
    struct callback_head     mce_kill_me;
    int                     mce_count;
#endif

#ifdef CONFIG_KRETPROBES
    struct llist_head        kretprobe_instances;
#endif
#ifdef CONFIG_RETHOOK
    struct llist_head        rethooks;
#endif

#ifdef CONFIG_ARCH_HAS_PARANOID_L1D_FLUSH
    /*
     * If L1D flush is supported on mm context switch
     * then we use this callback head to queue kill work
     * to kill tasks that are not running on SMT disabled
     * cores
     */
    struct callback_head     l1d_flush_kill;
#endif

#ifdef CONFIG_RV
    /*
     * Per-task RV monitor. Nowadays fixed in RV_PER_TASK_MONITORS.
     * If we find justification for more monitors, we can think
     * about adding more or developing a dynamic method. So far,
     * none of these are justified.
     */
    union rv_task_monitor    rv[RV_PER_TASK_MONITORS];
#endif

#ifdef CONFIG_USER_EVENTS
    struct user_event_mm     *user_event_mm;
#endif

    /*
     * New fields for task_struct should be added above here, so that
     * they are included in the randomized portion of task_struct.
     */

```

```

    randomized_struct_fields_end

    /* CPU-specific state of this task: */
    struct thread_struct      thread;

    /*
     * WARNING: on x86, 'thread_struct' contains a variable-sized
     * structure.  It *MUST* be at the end of 'task_struct'.
     *
     * Do not put anything below here!
     */
};

```

include/linux/sched/signal.h

```

#define tasklist_empty() \
    list_empty(&init_task.tasks)

#define next_task(p) \
    list_entry_rcu((p)→tasks.next, struct task_struct, tasks)

#define for_each_process(p) \
    for (p = &init_task ; (p = next_task(p)) ≠ &init_task ; )

```

End.