

DATABASE SYSTEM PRINCIPLE - TRANSACTION OVERVIEW

李旭东 Li-Xudong

LEEXUDONG@NANKAI.EDU.CN
NANKAI UNIVERSITY

OBJECTIVES

- Transaction Concept
- ACID Properties
- How to ensure ACID
- Multi-Transactions Failures: Isolation and Atomicity
- Multi-Transactions Isolation Levels
- How to provide concurrency schedules
- Transactions as SQL Statement

©LXD


OBJECTIVES

- **Transaction Concept**
- ACID Properties
- How to ensure ACID
- Multi-Transactions Failures: Isolation and Atomicity
- Multi-Transactions Isolation Levels
- How to provide concurrency schedules
- Transactions as SQL Statement

©LXD

APP: PROGRAM

```
APP
read (A)
A := A - 10
write (A)
read (B)
B := B + 10
write (B)
```



©LXD

A SIMPLE BUSINESS (1.1)

```
T0
read (A)
A := A - 10
write (A)
read (B)
B := B + 10
write (B)
```

- A business(collections of operations): transfer \$10 from account A to account B
 - Initial value: A=1000, B=2000
 - ?result
 - If a failure happens in read(B), ...?
 - In fact, We hope that None occur!!!

©LXD

A SIMPLE BUSINESS (1.2)

```
T0
read (A)
A := A - 10
write (A)
read (B)
B := B + 10
write (B)
```

- A business(collections of operations): transfer \$10 from account A to account B
 - (1) Atomicity: as a single unit, all or none
 - (2) Consistency:
 - consistent state(A+B unchanged)
 - (3) Durability:
 - if execution completes successfully, all the updates persist, even if a system failure

©LXD

A SIMPLE BUSINESS (2.1)

T_0	T_1
<pre> read (A) A := A - 10 write (A) read (B) B := B + 10 write (B) </pre>	<pre> read (A) A := A - 50 write (A) read (B) B := B + 50 write (B) </pre>

- There are two businesses:
 - T_0 : transfer \$10 from account A to account B
 - T_1 : transfer \$50 from account A to account B
- Initial value: A=1000, B=2000
- Interleaved交叉 (Concurrency): ?result
- Sequential(T_0, T_1): ?result
- Sequential(T_1, T_0): ?result
- ? Which is correct result ?

Sequential
Execution:
OK

©LXD

A SIMPLE BUSINESS (2.1)

T_0	T_1
<pre> read (A) A := A - 10 write (A) read (B) B := B + 10 write (B) </pre>	<pre> read (A) A := A - 50 write (A) read (B) B := B + 50 write (B) </pre>

T_0	T_1
<pre> read (A) A := A - 10 write (A) read (B) B := B + 10 write (B) </pre>	<pre> read (A) A := A - 50 write (A) read (B) B := B + 50 write (B) </pre>

- A business: transfer \$10 from account A to account B
- (4) Isolation:
 - even though multiple businesses may execute concurrently, the system guarantees that every business executes in turn (Sequentially, serially)

©LXD

A SIMPLE BUSINESS (3)

T_0
<pre> Begin transaction read (A) A := A - 10 write (A) read (B) B := B + 10 write (B) End transaction </pre>

- We need to invent **Transaction** concept
- collections of operations that form a single logical unit of work which do all or none.
- Begin transaction, ..., end transaction
- ACID Properties of Transaction
 - (1) Atomicity: as a single unit, all or none
 - (2) Consistency: consistent state(A+B unchanged)
 - (3) Durability: if execution completes successfully, all the updates persist, even if a system failure
 - (4) Isolation: even though multiple businesses may execute concurrently, the system guarantees that every business executes in turn (Sequentially, serially)

©LXD

OBJECTIVES

- Transaction Concept
- ACID Properties
- How to ensure ACID
- Multi-Transactions Failures: Isolation and Atomicity
- Multi-Transactions Isolation Levels
- How to provide concurrency schedules
- Transactions as SQL Statement

©LXD

TRANSACTION: ACID PROPERTIES

A **transaction** is a unit of program execution that accesses and possibly updates various data items. To preserve the integrity of data the database system must ensure:

- **Atomicity**. Either all operations of the transaction are properly reflected in the database or none are.
- **Consistency**. Execution of a transaction in isolation preserves the consistency of the database.
- **Isolation**. Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions. Intermediate transaction results must be hidden from other concurrently executed transactions.
 - That is, for every pair of transactions T_i and T_j , it appears to T_i that either T_j finished execution before T_i started, or T_j started execution after T_i finished.
- **Durability**. After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.

©LXD

OBJECTIVES

- Transaction Concept
- ACID Properties
- How to ensure ACID
- Multi-Transactions Failures: Isolation and Atomicity
- Multi-Transactions Isolation Levels
- How to provide concurrency schedules
- Transactions as SQL Statement

©LXD

TRANSACTION: HOW TO ENSURE THE ACID

- (1)How to ensure the consistency?
 - (1.1)for a transaction, the logical operations must be correct
 - Erroneous transaction logic can lead to inconsistency
 - (1.2)If we can ensure the atomicity, durability and isolation
 - (1.3)*Don't worry that during transaction execution the database may be temporarily inconsistent*

©LXD

TRANSACTION: HOW TO ENSURE THE ACID

- (2)How to ensure the durability?
 - Stable Storage
 - (transaction) Log

©LXD

TRANSACTION: HOW TO ENSURE THE ACID

- Storage Hierarchy
 - Volatile storage、Nonvolatile Storage

Typical access time		Typical capacity
1 nsec	Registers	<1 KB
2 nsec	Cache	4 MB
10 nsec	Main memory	512-2048 MB
10 msec	Magnetic disk	200-1000 GB
100 sec	Magnetic tape	400-800 GB

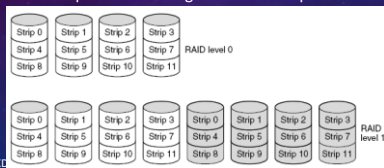
TRANSACTION: HOW TO ENSURE THE ACID

- Stable Storage
 - Information residing in stable storage is **never** lost!
 - Theoretically never cannot be guaranteed!
 - But, for any **a single nonvolatile storage** (magnetic disk or flash storage) , it is not absolutely reliable or safe!
 - How to implement **stable storage**?

©LXD

TRANSACTION: HOW TO ENSURE THE ACID

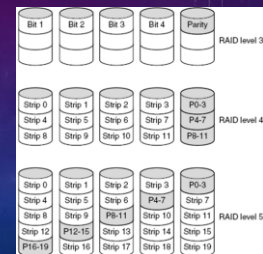
- Stable Storage **candidate**
 - RAID: Redundant Array of Independent Disks
 - To Improve disk storage safe and disk performance



©LXD

TRANSACTION: HOW TO ENSURE THE ACID

- Stable Storage **candidate**
 - RAID



©LXD

TRANSACTION: HOW TO ENSURE THE ACID

• Stable Storage **candidates**

- RAID
- ?LVM
- ?Cloud
- ...

©LXD

TRANSACTION: HOW TO ENSURE THE ACID

• (2)How to ensure the durability?

- Stable Storage
- (transaction) Log
 - for update log record
 - Transaction identifier、Data-item identifier
 - Old value, New value

©LXD

TRANSACTION: LOG

- A **log** is kept on stable storage: a sequence of **log records**
- When transaction T_i starts, it registers itself by writing a record $\langle T_i, \text{start} \rangle$ to the log
- Before T_i executes **write**(X), a log record $\langle T_i, X, V_1, V_2 \rangle$ is written, where V_1 is the value of X before the write (the **old value**), and V_2 is the value to be written to X (the **new value**).
- When T_i finishes its last statement, the log record $\langle T_i, \text{commit} \rangle$ is written.
- Two approaches using logs
 - (1) Immediate database modification
 - (2) Deferred database modification

©LXD

TRANSACTION: HOW TO ENSURE THE ACID

• (3)How to ensure the atomicity?

- Aborted transaction
 - The transaction doesn't complete its execution successfully
 - **Rolled back by log**
- Committed transaction
 - The transaction completes its execution successfully

©LXD

TRANSACTION: HOW TO ENSURE THE ACID

• (3)How to ensure the atomicity?

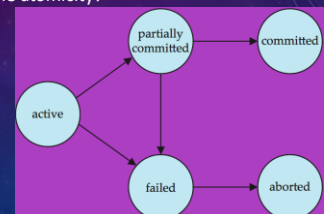
- **Compensating transaction** 补偿事务
 - Once a transaction has committed, we cannot undo its effects by aborting it.
 - The only way to undo the effects of a committed transaction is execute another transaction which do the opposite operations
- example
 - A transaction : transfer \$10 from account A to account B
 - A **compensating transaction** : transfer \$10 from account B to account A

©LXD

TRANSACTION: HOW TO ENSURE THE ACID

• (3)How to ensure the atomicity?

- Transaction model
 - Transaction states



©LXD

TRANSACTION STATES

- **Active** – the initial state; the transaction stays in this state while it is executing
- **Partially committed** – after the final statement(code) has been executed.
- **Failed** -after the discovery that normal execution can no longer proceed.
- **Aborted** – after the transaction has been rolled back and the database restored to its state prior to the start of the transaction.
 - Two options after it has been aborted:
 - Restart the transaction : can be done only if no internal logical error
 - Kill the transaction
- **Committed** – after successful completion.

©LXD

TRANSACTION: HOW TO ENSURE THE ACID

- (3)How to ensure the atomicity?
 - Use **transaction**

T_0	T_0	T_0	T_0
Begin transaction read (A) $A := A - 10$ write (A) read (B) rollback $B := B + 10$ write (B) End transaction	Begin transaction read (A) $A := A - 10$ write (A) read (B) rollback $B := B + 10$ write (B) End transaction	Begin transaction read (A) $A := A - 10$ write (A) read (B) $B := B + 10$ write (B) rollback End transaction	Begin transaction read (A) $A := A - 10$ write (A) read (B) $B := B + 10$ write (B) commit End transaction

©LXD

TRANSACTION: HOW TO ENSURE THE ACID

- (4)How to ensure the isolation?
 - Isolation in only one transaction in whole system
 - No problem
 - Isolation involved multi transactions
 - even though multiple businesses may execute concurrently, the system guarantees that every business executes in turn (Sequentially, serially)

©LXD

TRANSACTION: HOW TO ENSURE THE ISOLATION INVOLVED MULTI TRANSACTIONS?

T_0
read (A) $A := A - 10$ write (A) read (B) $B := B + 10$ write (B)

T_1
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B)

? What is the **chronological order** in which instructions of **concurrent** transactions are executed?

©LXD

TRANSACTION: HOW TO ENSURE THE ISOLATION INVOLVED MULTI TRANSACTIONS?

Case 1		Case 2		Case N	
T_0	T_1	T_0	T_1	T_0	T_1
read (A) $A := A - 10$ write (A) read (B) $B := B + 10$ write (B)	read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B)	read (A) $A := A - 10$ write (A) read (B) $B := B + 10$ write (B)	read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B)	read (A) $A := A - 10$ write (A) read (B) $B := B + 10$ write (B)	read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B)

©LXD

In different cases, time sequence of The CPU Instructors is different

TRANSACTION SCHEDULER: SCHEDULE

- **Schedule** – a **sequences of instructions** that specify the chronological order in which instructions of concurrent transactions are executed
 - A schedule for a set of transactions must consist of all instructions of those transactions
 - Must preserve the order in which the instructions appear in each individual transaction.
- A transaction that successfully completes its execution will have a **commit** instructions as the last statement
 - By default transaction assumed to execute commit instruction as its last step
- A transaction that fails to successfully complete its execution will have an **abort** instruction as the last statement

©LXD

TRANSACTION: HOW TO ENSURE THE ISOLATION INVOLVED MULTI TRANSACTIONS?

Sa1:

T_1	T_2
<pre> read (A) A := A - 50 write (A) read (B) B := B + 50 write (B) </pre>	<pre> read (A) temp := A * 0.1 A := A - temp write (A) read (B) B := B + temp write (B) </pre>

- There are two truncations:
 - T1: transfer \$50 from account A to account B
 - T2: transfer A*0.1 from account A to account B
- Initial value: A=1000, B=2000
- ?result **OK**
- ? Does it satisfy the consistency?

TRANSACTION: HOW TO ENSURE THE ISOLATION INVOLVED MULTI TRANSACTIONS?

Sa2:

T_1	T_2
<pre> read (A) A := A - 50 write (A) read (B) B := B + 50 write (B) </pre>	<pre> read (A) temp := A * 0.1 A := A - temp write (A) read (B) B := B + temp write (B) </pre>

- There are two truncations:
 - T1: transfer \$50 from account A to account B
 - T2: transfer A*0.1 from account A to account B
- Initial value: A=1000, B=2000
- ?result **OK**
- ? Does it satisfy the consistency?

TRANSACTION: HOW TO ENSURE THE ISOLATION INVOLVED MULTI TRANSACTIONS?

Sa3:

T_1	T_2
<pre> read (A) A := A - 50 write (A) read (B) B := B + 50 write (B) </pre>	<pre> read (A) temp := A * 0.1 A := A - temp write (A) read (B) B := B + temp write (B) </pre>

- There are two truncations:
 - T1: transfer \$50 from account A to account B
 - T2: transfer A*0.1 from account A to account B
- Initial value: A=1000, B=2000
- ?result **OK**
- ? Does it satisfy the consistency?

TRANSACTION: HOW TO ENSURE THE ISOLATION INVOLVED MULTI TRANSACTIONS?

Sa4:

T_1	T_2
<pre> read (A) A := A - 50 write (A) read (B) B := B + 50 write (B) </pre>	<pre> read (A) temp := A * 0.1 A := A - temp write (A) read (B) B := B + temp write (B) </pre>

- There are two truncations:
 - T1: transfer \$50 from account A to account B
 - T2: transfer A*0.1 from account A to account B
- Initial value: A=1000, B=2000
- ?result **error**
- ? Does it satisfy the consistency?

TRANSACTION: HOW TO ENSURE THE ISOLATION INVOLVED MULTI TRANSACTIONS?

- How to ensure the isolation with multi-transaction
 - **The first solution:** Insist all the related transactions run serially
- Some Drawbacks of serial executions
 - Long waiting time
 - Run these transactions slowly
 - Low throughput
 - In given amount of time, the system just do little transactions
 - Low resource utilization

©LXD • CPU, I/O, Disk, Memory

TRANSACTION: HOW TO ENSURE THE ISOLATION INVOLVED MULTI TRANSACTIONS?

- How to ensure the isolation with multi-transaction
 - **The second solution:** concurrency并发, but conflict serializability冲突可串行化
 - **Concurrent executions**
 - **Serializability Theory**
 -

©LXD

CONCURRENT EXECUTIONS 并发执行


- Multiple transactions are allowed to run **concurrently** in the system. Advantages are:
 - Increased processor and disk utilization**, leading to better transaction **throughput**
 - E.g. one transaction can be using the CPU while another is reading from or writing to the disk
 - Reduced average response time** for transactions: short transactions need not wait behind long ones.
- Concurrency control schemes** – mechanisms to achieve isolation
 - That is, to control the interaction among the concurrent transactions in order to prevent them from destroying the consistency of the database

©LXD

TRANSACTION: SERIALIZABILITY THEORY

- Abstract database operations

- Read(Q), Write(Q)



©LXD

TRANSACTION: SERIALIZABILITY THEORY

- For a schedule (s) of two transactions

- Example

- there two consecutive instructions, I and J
- $I \in T_i, J \in T_j$

T_i	T_j
I	J
...	...

©LXD

TRANSACTION: SERIALIZABILITY THEORY

- For a schedule (s) of two transactions (cont.,)

- Example

- Case 1: If I and J refer to different data, then we can swap I and J without affecting the results of any instruction in the schedule s

T_i	T_j
I: Write(A)	J: Write(B)
...	...

©LXD

TRANSACTION: SERIALIZABILITY THEORY

- For a schedule (s) of two transactions (cont.,)

- Example

- Case 2: If I and J refer to the same data, then the order of the two steps may **matter** ...

T_i	T_j
I: Write(A)	J: Write(A)
...	...

T_j	T_i
I: Write(A)	J: Write(A)
...	...

Cannot swap I and J

©LXD

TRANSACTION: SERIALIZABILITY THEORY

- For a schedule (s) of two transactions (cont.,)

- Case 2: If I and J refer to the same data, then the order of the two steps may **matter** ...

Only the first case, We can swap I and J, regardless of the order

(1) I=Read(A), J=Read(A)	(2) I=Read(A), J=Write(A)
(3) I=Write(A), J=Write(A)	(4) I=Write(A), J=Read(A)

Conflict

- We say that I and J conflict
- If they are operations by **different transactions** on the **same data** item,
- And at least one of these instructions is a **write** operation.

©LXD


CONFLICTING INSTRUCTIONS

- Let I_i and I_j be two Instructions of transactions T_i and T_j respectively. Instructions I_i and I_j **conflict** if and only if there exists some item Q accessed by both I_i and I_j , and at least one of these instructions wrote Q .
 - $I_i = \text{read}(Q)$, $I_j = \text{read}(Q)$. I_i and I_j don't conflict.
 - $I_i = \text{read}(Q)$, $I_j = \text{write}(Q)$. They conflict.
 - $I_i = \text{write}(Q)$, $I_j = \text{read}(Q)$. They conflict.
 - $I_i = \text{write}(Q)$, $I_j = \text{write}(Q)$. They conflict.
- Intuitively, a conflict between I_i and I_j forces a (logical) temporal order between them.
 - If I_i and I_j are consecutive in a schedule and they do not conflict, their results would remain the same even if they had been interchanged in the schedule.

©LYD

TRANSACTION: SERIALIZABILITY THEORY

- conflict serializable**




- So, S1 and S2 and S3 are conflict equivalent
- So, S1 and S2 are both conflict serializable

©LYD

TRANSACTION: SERIALIZABILITY THEORY

- conflict serializable**



- S4 and S5 are not conflict equivalent

©LYD

TRANSACTION: SERIALIZABILITY THEORY

- How to judge whether or not a set of transaction are **conflict serializable**?

- Precedence graph** (前驱图、优先图)

- $G=(V,E)$

- V : a set of vertices (transaction names)

- E : a set of edges (conflict operation)

- $T_i \rightarrow T_j$: three conditions holds

- (1) T_i executes write(Q) before T_j executes read(Q)

- (2) T_i executes read(Q) before T_j executes write(Q)


- (1) T_i executes write(Q) before T_j executes write(Q)

©LYD

PRECEDENCE GRAPH

- Consider some schedule of a set of transactions T_1, T_2, \dots, T_n
- Precedence graph** — a direct graph where the vertices are the transactions (names).
- We draw an arc from T_i to T_j ($T_i \rightarrow T_j$) if the two transaction conflict, and T_i accessed the data item on which the conflict arose earlier.
- We may label the arc by the item that was accessed.

- Example**



©LYD

TESTING FOR CONFLICT SERIALIZABILITY

- A schedule is **conflict serializable** if and only if its precedence graph is **acyclic** (非循环).
- Cycle-detection algorithms** exist which take order n^2 time, where n is the number of vertices in the graph.
 - (Better algorithms take order $n + e$ where e is the number of edges.)
- If precedence graph is acyclic, the **serializability order** can be obtained by a **topological sorting** of the graph.
 - That is, a linear order consistent with the partial order of the graph.
 - For example, a serializability order for the schedule (a) would be one of either (b) or (c)



TRANSACTION: SERIALIZABILITY THEORY

- An **exception** ^{S6} to between Precedence graph and conflict serializable ^{S7}

T_1	T_2
read (A) $A := A - 50$ write (A)	
	read (B) $B := B - 10$ write (B)
read (B) $B := B + 50$ write (B)	
	read (A) $A := A + 10$ write (A)

- S6 precedence graph is acyclic?
- Is S6 conflict equivalent to S7?
- The outcomes of S6 and S7?

view equivalence
视图等价

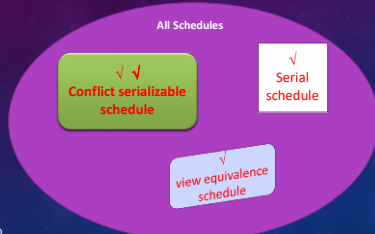
T_1	T_2
read (A) $A := A - 50$ write (A)	
	read (B) $B := B + 50$ write (B)
	read (B) $B := B - 10$ write (B)
	read (A) $A := A + 10$ write (A)

SUMMARY: HOW TO ENSURE THE ISOLATION INVOLVED MULTI TRANSACTIONS? (1/2)

- How to ensure the isolation with multi-transactions

- The first solution:**
 - Insist all the related transactions run serially
- The second solution:**
 - concurrency 并发, but conflict serializability 冲突可串行化

SUMMARY: HOW TO ENSURE THE ISOLATION INVOLVED MULTI TRANSACTIONS? (2/2)



OBJECTIVES

- Transaction Concept
- ACID Properties
- How to ensure ACID
- Multi-Transactions Failures: Isolation and Atomicity**
- Multi-Transactions Isolation Levels
- How to provide concurrency schedules
- Transactions as SQL Statement

MULTI-TRANSACTIONS FAILURES:
- ISOLATION AND ATOMICITY

- For Multi-Transactions
 - Case 1: these transactions are mutually independent
 - If one transaction failed, just rollback this transaction
 - Case 2: these transactions are mutually dependent
 - T_j is dependent on T_i : T_j has read data written by T_i
 - If T_i transaction failed, how to rollback?

RECOVERABLE SCHEDULES 可恢复调度

- Recoverable schedule** — if a transaction T_j reads a data item previously written by a transaction T_i , then the commit operation of T_i must appear before the commit operation of T_j .
- The following schedule is not recoverable if T_2 commits immediately after the read(A) operation.

T_1	T_2
read (A) write (A)	
	read (A) commit
read (B)	

S8 is a nonrecoverable schedule

- If T_1 should abort, T_2 would have read (and possibly shown to the user) an inconsistent database state. Hence, database must ensure that schedules are recoverable.

CASCADING ROLLBACKS级联回滚

- **Cascading rollback** — a single transaction failure leads to a series of transaction rollbacks.
- Consider the following schedule where none of the transactions has yet committed (so the schedule is **recoverable**)

S9

T_8	T_9	T_{10}
read(A) read(B) write(A)		
abort	read(A) write(A)	read(A)

If T_8 fails, T_9 and T_{10} must also be rolled back.

- Can lead to the undoing of a significant amount of work

©LXD

CASCADELESS SCHEDULES无级联调度

- **Cascadeless schedules** — for each pair of transactions T_i and T_j such that T_i reads a data item previously written by T_j , the **commit operation** of T_j appears before the read operation of T_i .
- Every cascadeless schedule is also recoverable
- It is desirable to restrict the schedules to those that are cascadeless
- Example of a schedule that is **NOT** cascadeless

S10

T_8	T_9	T_{10}
read(A) read(B) write(A) abort		
	read(A) write(A)	
		read(A)

- When T8 is rolled back
 - T9 and T10 don't need to be rolled back

©LXD

OBJECTIVES

- Transaction Concept
- ACID Properties
- How to ensure ACID
- Multi-Transactions Failures: Isolation and Atomicity
- **Multi-Transactions Isolation Levels**
- How to provide concurrency schedules
- Transactions as SQL Statement

©LXD

CONSISTENCY (1/2)

- 1. **Strong Consistency: serializable**
 - All accesses are seen by all parallel processes (or nodes, processors etc.) in the same order (sequentially)
 - Serializability is a useful concept because it allows programmers to ignore issues related to concurrency when they code transactions

©LXD

CONSISTENCY (2/2)

- 2. **Weak Consistency: not serializable**
 - (1) All accesses to **synchronization variables** are seen by all processes (or nodes, processors) in the same order (sequentially) - these are synchronization operations. Accesses to critical sections are seen sequentially.
 - (2) All other accesses may be seen in different order on different processes (or nodes, processors).
 - (3) The set of both read and write operations in between different synchronization operations is the same in each process.

©LXD

WEAK LEVELS OF CONSISTENCY

- Some applications are willing to live with weak levels of consistency, allowing schedules that are **not serializable**
 - E.g., a read-only transaction that wants to get an approximate total balance of all accounts
 - E.g., database statistics computed for query optimization can be approximate (why?)
 - Such transactions need not be serializable with respect to other transactions
- **Tradeoff** accuracy for performance

©LXD

LEVELS OF CONSISTENCY IN SQL-92

- **Serializable** — default
- **Repeatable read** — only committed records to be read, repeated reads of same record must return same value. However, a transaction may not be serializable — it may find some records inserted by a transaction but not find others.
- **Read committed** — only committed records can be read, but successive reads of record may return different (but committed) values.
- **Read uncommitted** — even uncommitted records may be read.
- **SQL cmd:** set transaction isolation level serializable;

©LXD

LEVELS OF CONSISTENCY

- **Dirty writes** 脏写
 - All the isolation levels above additionally **disallow** dirty writes
 - That is, the disallow writes to a data item that has already been written by another transaction that has not yet committed or aborted

©LXD

OBJECTIVES

- Transaction Concept
- ACID Properties
- How to ensure ACID
- Multi-Transactions Failures: Isolation and Atomicity
- Multi-Transactions Isolation Levels
- **How to provide concurrency schedules**
- Transactions as SQL Statement

©LXD

HOW TO PROVIDE CONCURRENCY SCHEDULES

- A database must provide a mechanism that will ensure that all possible schedules are:
 - **Conflict serializable** (or View serializable)
 - **Recoverable**
 - **preferably cascadeless** 无级联
- (1) A policy in which only one transaction can execute at a time generates **serial schedules**, but provides a poor degree of concurrency
- ©LXD

GOAL OF CONCURRENCY CONTROL

- (2) Concurrency-control schemes **tradeoff**
 - between the amount of concurrency they allow and the amount of overhead that they incur
- Testing a schedule for serializability **after** it has executed is a little too late!
 - Tests for serializability help us understand why a concurrency control protocol is correct
- **Goal** — to develop **concurrency control protocols** that will assure **serializability** and provide a high degree of concurrency.

©LXD

CONCURRENCY CONTROL PROTOCOLS

- **Locking** (in this ppt)
- Graph-based protocols
- Timestamps
- Validation-based protocols
- Multiple versions and snapshot isolation
- ...

©LXD

LOCK-BASED PROTOCOLS

- A lock is a mechanism to control concurrent access to a data item
- Data items can be locked in two modes :
 - exclusive (X) mode 排他锁** Data item can be both read as well as written. X-lock is requested using **lock-X** instruction.
 - shared (S) mode 共享锁** Data item can only be read. S-lock is requested using **lock-S** instruction.
- Lock requests are made to the concurrency-control manager by the programmer. Transaction can proceed only after request is **granted 授权**.

©LYD

LOCK-BASED PROTOCOLS (CONT.)

- Lock-compatibility matrix**

	S	X
S	true	false
X	false	false

- * A transaction may be **granted** a lock on an item if the requested lock is compatible with locks already held on the item by other transactions
- Any number of transactions can hold shared locks on an item,
 - But if any transaction holds an exclusive on the item no other transaction may hold any lock on the item.
- If a lock cannot be **granted**
 - the requesting transaction is made to **wait** till all incompatible locks held by other transactions have been released. The lock is then granted.

LOCK-BASED PROTOCOLS (CONT.)

- When a transaction T_i requests a lock on a data item Q in a particular mode M
 - The concurrency-control manager **grants** the lock provided that:
 - (1) There is no other transaction holding a lock on Q in a mode that conflicts with M
 - (2) There is no other transaction that is waiting for a lock on Q and that made its lock request before T_i

©LYD

EXAMPLE(A): LOCK-BASED PROTOCOLS

- SOLUTION 1

T_{11}	T_{12}
read (B) $B := B - 50$ write (B) read (A) $A := A + 50$ write (A)	read (B) read (A) Display(A+B)



T_{11}	T_{12}
Lock-X(B) read (B) $B := B - 50$ write (B) Unlock(B) Lock-X(A) read (A) $A := A + 50$ write (A) Unlock(A)	Lock-S(B) read (B) Unlock(B) Lock-S(A) read (A) Unlock(A) Display(A+B)

©LYD

EXAMPLE(A):
LOCK-BASED PROTOCOLS
- SOLUTION 1

T_{11}	T_{12}
read (B) $B := B - 50$ write (B) read (A) $A := A + 50$ write (A)	read (B) read (A) Display(A+B)

Initial values:
 $A=100$
 $B=200$

error

©LYD

EXAMPLE(A): LOCK-BASED PROTOCOLS

- SOLUTION 2

T_{11}	T_{12}
read (B) $B := B - 50$ write (B) read (A) $A := A + 50$ write (A)	read (B) read (A) Display(A+B)



T_{11}	T_{12}
Lock-X(B) read (B) $B := B - 50$ write (B) Lock-X(A) read (A) $A := A + 50$ write (A) Unlock(A) Unlock(B)	Lock-S(B) read (B) Lock-S(A) read (A) Display(A+B) Unlock(A) Unlock(B)

OK!

©LYD

EXAMPLE(A): LOCK-BASED PROTOCOLS - SOLUTION 2

The Two-Phase Locking Protocol

T_{11}	T_{12}
read (B) $B := B - 50$ write (B)	
read (A) $A := A + 50$ write (A)	

Initial values:
A=100
B=200



T_{11}	T_{12}	Concurrency-control manager
Lock-x(B)		Grant-x(B, T_{11})
read (B) $B := B - 50$ write (B)		
	Lock-S(B)	Grant-S(B, T_{12})
	read (B) Lock-S(A)	Grant-S(A, T_{12})
	read (A) Display(A+B) Unlock(A) Unlock(B)	
Lock-x(A)		Grant-x(A, T_{11})
read (A) $A := A + 50$ write (A) Unlock(A) Unlock(B)		

THE TWO-PHASE LOCKING PROTOCOL

- This protocol ensures conflict-serializable schedules.
- Phase 1: Growing Phase** 增长阶段
 - Transaction may obtain locks
 - Transaction may not release locks
- Phase 2: Shrinking Phase** 缩减阶段
 - Transaction may release locks
 - Transaction may not obtain locks
- The protocol assures **serializability**. It can be proved that the transactions can be serialized in the order of their **lock points** 锁点 (i.e., the point where a transaction acquired its final lock).

EXAMPLE(1): THE TWO-PHASE LOCKING PROTOCOL

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

- There are two truncations:
 - T1: transfer \$50 from account A to account B
 - T2: transfer $A * 0.1$ from account A to account B
- Initial value: A=1000, B=2000

EXAMPLE(1): THE TWO-PHASE LOCKING PROTOCOL

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit



T_1	T_2
Lock-x(A) read (A) $A := A - 50$ write (A) Lock-x(B) read (B) $B := B + 50$ write (B) Unlock(B) Unlock(A)	Lock-x(A) read (A) $temp := A * 0.1$ $A := A - temp$ write (A) Lock-x(B) read (B) $B := B + temp$ write (B) Unlock(B) Unlock(A)

EXAMPLE(1): THE TWO-PHASE LOCKING PROTOCOL

Sa1:

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B)	
	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B)

T_1	T_2
Lock-x(A) read (A) $A := A - 50$ write (A) Lock-x(B) read (B) $B := B + 50$ write (B) Unlock(B) Unlock(A)	
	Lock-x(A) read (A) $temp := A * 0.1$ $A := A - temp$ write (A) Lock-x(B) read (B) $B := B + temp$ write (B) Unlock(B) Unlock(A)

EXAMPLE(1): THE TWO-PHASE LOCKING PROTOCOL

Sa2:

T_1	T_2
	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B)
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B)	

T_1	T_2
	Lock-x(A) read (A) $temp := A * 0.1$ $A := A - temp$ write (A) Lock-x(B) read (B) $B := B + temp$ write (B) Unlock(B) Unlock(A)
Lock-x(A) read (A) $A := A - 50$ write (A) Lock-x(B) read (B) $B := B + 50$ write (B) Unlock(B) Unlock(A)	

EXAMPLE(1): THE TWO-PHASE LOCKING PROTOCOL

Sa3:

T_1	T_2
read (A) $A := A - 50$ write (A)	read (A) $temp := A * 0.1$ $A := A - temp$ write (A)
read (B) $B := B + 50$ write (B)	read (B) $B := B + temp$ write (B)

©LYD

T_1	T_2
Lock-x(A) read (A) $A := A - 50$ write (A)	Lock-x(A) read (A) $temp := A * 0.1$ $A := A - temp$ write (A)
Lock-x(B) read (B) $B := B + 50$ write (B) Unlock(B) Unlock(A)	Lock-x(B) read (B) $B := B + temp$ write (B) Unlock(B) Unlock(A)

EXAMPLE(1): THE TWO-PHASE LOCKING PROTOCOL

Sa4:

T_1	T_2
read (A) $A := A - 50$ write (A)	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B)
read (B) $B := B + 50$ write (B)	read (B) $B := B + temp$ write (B)

©LYD

Sa4 is not conflict serializable

T_1	T_2
Lock-x(A) read (A) $A := A - 50$ write (A)	Lock-x(A) read (A) $temp := A * 0.1$ $A := A - temp$ write (A) Lock-x(B) read (B)
Lock-x(B) read (B) $B := B + 50$ write (B) Unlock(B) Unlock(A)	Lock-x(B) read (B) $B := B + temp$ write (B) Unlock(B) Unlock(A)

EXAMPLE(2): THE TWO-PHASE LOCKING PROTOCOL (1/2)

T_{11}	T_{12}
read (B) $B := B - 50$ write (B) read (A) $A := A + 50$ write (A)	read (A) read (B) Display(A+B)



T_{11}	T_{12}
Lock-X(B) read (B) $B := B - 50$ write (B) Lock-X(A) read (A) $A := A + 50$ write (A) Unlock(A) Unlock(B)	Lock-S(A) read (A) Lock-S(B) read (B) Display(A+B) Unlock(A)

©LYD

EXAMPLE(2): THE TWO-PHASE LOCKING PROTOCOL (2/2)

T_{11}	T_{12}
read (B) $B := B - 50$ write (B) read (A) $A := A + 50$ write (A)	read (A) read (B) Display(A+B)

Deadlock!

T_{11}	T_{12}	Concurrency-controlled manager
Lock-x(B) read (B) $B := B - 50$ write (B)		Grant-x(B, T ₁₁)
	Lock-S(A)	Grant-S(A, T ₁₂)
	read (A) Lock-S(B)	Grant-S(B, T ₁₂)
	read (B) Display(A+B) Unlock(B) Unlock(A)	
Lock-X(A) read (A) $A := A + 50$ write (A) Unlock(A) Unlock(B)		Grant-x(A, T ₁₁)

©LYD

EXAMPLE(3): THE TWO-PHASE LOCKING PROTOCOL

S6:

T_1	T_2
read (A) $A := A - 50$ write (A)	read (B) $B := B - 10$ write (B)
read (B) $B := B + 50$ write (B)	read (A) $A := A + 10$ write (A)

1. How to follow two-phase locking?
2. The result of schedule?

©LYD

THE TWO-PHASE LOCKING PROTOCOL

- Two-phase locking **does not** ensure freedom from **deadlock**

©LYD

THE TWO-PHASE LOCKING PROTOCOL

- Variant of Two-phase locking
 - **Strict two-phase locking protocol**
 - All exclusive-mode locks taken by a transaction be held until that transaction commits
 - **Rigorous two-phase locking protocol**
 - All locks be held by a transaction be held until that transaction commits

©LXD

THE TWO-PHASE LOCKING PROTOCOL

- Lock conversions转化 of Two-phase locking
 - **upgrade**
 - From shared mode lock to exclusive mode
 - **downgrade**
 - From exclusive mode lock to shared mode

©LXD

THE TWO-PHASE LOCKING PROTOCOL


- There can be **conflict serializable** schedules that cannot be obtained if two-phase locking is used.
- However, in the absence of extra information (e.g., ordering of access to data), two-phase locking is needed for conflict serializability in the following sense:
 - Given a transaction T_i that does not follow two-phase locking, we can find a transaction T_j that uses two-phase locking, and a schedule for T_i and T_j that is not conflict serializable.

©LXD

IMPLEMENTATION OF LOCKING

- Lock table
 - Data items
 - I4, I7, I23, I44, I912
 - Granted locks
 - T23, T1, T8
 - Waiting locks
 - T23, T2

©LXD



OBJECTIVES

- Transaction Concept
- ACID Properties
- How to ensure ACID
- Multi-Transactions Failures: Isolation and Atomicity
- Multi-Transactions Isolation Levels
- How to provide concurrency schedules
- **Transactions as SQL Statement**

©LXD

TRANSACTIONS AS SQL STATEMENT

- Data manipulation language must include a construct for specifying the set of actions that comprise a transaction.
- In SQL, a transaction begins implicitly.
- A transaction in SQL ends by:
 - **Commit work** commits current transaction and begins a new one.
 - **Rollback work** causes current transaction to abort.
- In almost all database systems, by default, every SQL statement also commits implicitly if it executes successfully
 - Implicit commit can be turned off by a database directive
 - E.g. in JDBC, `connection.setAutoCommit(false);`

©LXD

```
#!/usr/bin/python3
import sys
import mysql.connector
config = {
    'host': 'localhost',
    'user': 'myuser',
    'password': 'xxx',
    'port': 3306,
    'database': 'dbsc1ab2018',
    'charset': 'utf8'
}

#db connect
try:
    cnn = mysql.connector.connect(**config)
except mysql.connector.Error as e:
    cnn = None
    print('connect fails!{}'.format(e))
#
if None==cnn:
    sys.exit()
#open cursor
cursor = cnn.cursor()
```

©LXD

```
try:
    cursor.execute('select name,salary from instructor01')
    for name,salary in cursor:
        print(name.decode('utf-8'),' ',salary)
#
    cursor.execute('update instructor01 set salary=salary-100 where name="Bondi"')
    cursor.execute('update instructor01 set salary=salary+100 where name="Dale"')
    cursor.execute('select name,salary from instructor01')
    for name,salary in cursor:
        print(name.decode('utf-8'),' ',salary)
#
    cnn.rollback()
```

©LXD

```
#
cursor.execute('select name,salary from instructor01')
for name,salary in cursor:
    print(name.decode('utf-8'),' ',salary)
except mysql.connector.Error as e:
    print('query error!{}'.format(e))
finally:
    cursor.close()
    cnn.close()
#
print('exit')
```

©LXD

SUMMARY

- Transaction Concept
- ACID Properties
- How to ensure ACID
- Multi-Transactions Failures: Isolation and Atomicity
- Multi-Transactions Isolation Levels
- How to provide concurrency schedules
- Transactions as SQL Statement

©LXD

How to get a high degree of concurrency?

Q&A?

THANKS!

leexudong@nankai.edu.cn